

# Lecture 11: Computational Efficiency

LING 1340/2340: Data Science for Linguists  
Jevon Heath (but really, Na-Rae Han and Dan Zheng)

# Objectives

---

- ▶ Big data considerations
- ▶ Computational efficiency
  - ◆ Memory vs. processing time
  - ◆ Algorithmic complexity
  - ◆ Big O notation
- ▶ Word embeddings

# The Yelp Dataset Challenge

► <https://www.yelp.com/dataset/challenge>



## Download The Data

The links to download the data will be valid for **30 seconds**.

### JSON

[Download JSON](#)

3.6 gigabytes compressed  
8.69 gigabytes uncompressed

1 .tar.gz file compressed  
6 .json files uncompressed

For more information on the JSON dataset, visit the [main dataset documentation](#) page.

### Photos

[Download photos](#)

7.22 gigabytes compressed  
7.67 gigabytes uncompressed

1 .tar.gz file compressed  
1 .json file and 1 folder containing 200,000 photos

# Working with big data files

```
naraehan@login0:zfs2/ling1340-2019s/shared_data/yelp_dataset_13
[naraehan@login0 yelp_dataset_13]$ ls -lah
total 5.1G
drwxr-xr-x  2 naraehan ling1340-2019s  10 Mar 21 12:03 .
drwxr-xr-x  6 naraehan ling1340-2019s   6 Mar 21 12:05 ..
-rw-r--r--  1 naraehan ling1340-2019s 132M Mar 21 12:03 business.json
-rw-r--r--  1 naraehan ling1340-2019s 390M Mar 21 12:03 checkin.json
-rw-r--r--  1 naraehan ling1340-2019s  99K Mar 21 12:03 Dataset_Challenge_Dataset_Agreement.pdf
-rw-r--r--  1 naraehan ling1340-2019s  25M Mar 21 12:03 photo.json
-rw-r--r--  1 naraehan ling1340-2019s  5.0G Mar 21 12:03 review.json
-rw-r--r--  1 naraehan ling1340-2019s 234M Mar 21 12:03 tip.json
-rw-r--r--  1 naraehan ling1340-2019s  2.4G Mar 21 12:03 user.json
-rw-r--r--  1 naraehan ling1340-2019s 110K Mar 21 12:03 Yelp_Dataset_Challenge_Round_13.pdf
[naraehan@login0 yelp_dataset_13]$ wc -l review.json user.json
6685900 review.json
1637138 user.json
8323038 total
```

- ▶ Each file is in JSON format, and they are huge:
  - ◆ review.json is 5.0GB with 6.7 million records
  - ◆ user.json is 2.4GB with 1.6 million records
- ← Too big to open in most text editors (Notepad++ couldn't.)
- ← How to explore them?

# Command line exploration

```
naraehan@login0:/zfs2/ling1340-2019s/shared_data/yelp_dataset_13
[naraehan@login0 yelp_dataset_13]$ wc -l review.json user.json
6685900 review.json
1637138 user.json
8323038 total
[naraehan@login0 yelp_dataset_13]$ head -1 review.json
{"review_id":"Q1sbwvVQXV2734tPgoKj4Q","user_id":"hG7b0MtEbXx5QzbzE6C_VA","business_id":"ujmEBvifdJM6h6RLv4wQIg",
,"stars":1.0,"useful":6,"funny":1,"cool":0,"text":"Total bill for this horrible service? Over $8Gs. These cro
oks actually had the nerve to charge us $69 for 3 pills. I checked online the pills can be had for 19 cents EA
CH! Avoid Hospital ERs at all costs.","date":"2013-05-07 04:34:36"}
[naraehan@login0 yelp_dataset_13]$ grep -i 'scrumptious' review.json | wc -l
9223
[naraehan@login0 yelp_dataset_13]$ grep -i 'horrible' review.json | wc -l
149059
[naraehan@login0 yelp_dataset_13]$ grep -i 'horrible' review.json | cut -d, -f4 | head -3
"stars":1.0
"stars":1.0
"stars":1.0
[naraehan@login0 yelp_dataset_13]$ grep -i 'horrible' review.json | cut -d, -f4 | sort | uniq -c | sort -nr
101450 "stars":1.0
21380 "stars":2.0
11463 "stars":3.0
8436 "stars":5.0
6330 "stars":4.0
[naraehan@login0 yelp_dataset_13]$ |
```

# Opening + processing big files

- ▶ How much resource does it take to process review.json file (5.0GB)?

```
process_reviews.py - C:\Users\narae\Documents\Data_Science\dataset\process_reviews.py (3.5.3)
File Edit Format Run Options Window Help
import pandas as pd
import sys
from collections import Counter

filename = sys.argv[1]

df = pd.read_json(filename, lines=True, encoding='utf-8')

print(df.head(5))

wtoks = ' '.join(df['text']).split()
wfreq = Counter(wtoks)
print(wfreq.most_common(20))

|
```

There's 5 GB

Another ~4 GB

Not as big

This code is NOT  
memory-efficient.

After exceeding the  
8GB default  
memory allocation,  
the job gets killed.

# Memory consideration

- ▶ How much space needed for bigrams? Trigrams?

Good news! These  
are built as  
generator objects  
and take up almost  
zero space.

```
process_reviews2.py - C:/Users/narae/Documents/Data_Science/dataset/process_reviews2.py (3.5.3)
File Edit Format Run Options Window Help
import pandas as pd
import sys
from collections import Counter
import nltk

filename = sys.argv[1]

df = pd.read_json(filename, lines=True, encoding='utf-8')
print(df.head(5))

wtoks = ' '.join(df['text']).split()
bigrams = nltk.bigrams(wtoks)
trigrams = nltk.trigrams(wtoks)

bifreq = Counter(bigrams)
print(bifreq.most_common(20))

trifreq = Counter(trigrams)
print(trifreq.most_common(20))
```

But these  
frequency  
counter objects  
will take up  
space.

Generator type objects take up little memory space; meant to be used in a loop-like environment.

Content has been exhausted

```
>>> import nltk
>>> sent = 'Colorless green ideas sleep oh so very furiously'
>>> toks = sent.split()
>>> toks
['Colorless', 'green', 'ideas', 'sleep', 'oh', 'so', 'very', 'furiously']
>>> bigrams = nltk.bigrams(toks)
>>> bigrams
<generator object bigrams at 0x00000236371E2BF8>
>>> for b in bigrams:
>>>     print(b)

('Colorless', 'green')
('green', 'ideas')
('ideas', 'sleep')
('sleep', 'oh')
('oh', 'so')
('so', 'very')
('very', 'furiously')
>>> bigrams
<generator object bigrams at 0x00000236371E2BF8>
>>> list(bigrams)
[]
>>> bigrams = nltk.bigrams(toks)
>>> list(bigrams)
[('Colorless', 'green'), ('green', 'ideas'), ('ideas', 'sleep'), ('sleep', 'oh'), ('oh', 'so'), ('so', 'very'), ('very', 'furiously')]
>>>
```

Casting as list.  
If you store the returned list, it will take up memory space.



# File opening & closing methods

```
f = open('review.json')
lines = f.readlines()
for l in lines:
    if 'horrible' in l:
        print(l)
f.close()
```

```
lines = open('review.json').readlines()
for l in lines :
    if 'horrible' in l:
        print(l)
```

Python will  
close up this  
file handle.

Which methods  
are memory-  
efficient?

```
f = open('review.json')
for l in f:
    if 'horrible' in l:
        print(l)
f.close()
```

```
with open('review.json') as f:
    for l in f:
        if 'horrible' in l:
            print(l)
```

No need to close f later.  
Some folks swear by using  
**with**.

# Handling files in chunks

```
f = open('review.json')
lines1 = f.readlines(1000000000)
lines2 = f.readlines(1000000000)
lines3 = f.readlines(1000000000)
lines4 = f.readlines(1000000000)
lines5 = f.readlines()
f.close()
```

Optional # of bytes to read.  
(When used like this without a loop, offers no memory advantage.)

Generator object. Takes up zero space.

```
dfs = pd.read_json('review.json', lines=True, chunksize=10000, encoding='utf8')
wfreq = Counter()

for df in dfs:
    wtoks = ' '.join(df['text']).split()
    temp = Counter(wtoks)
    wfreq.update(temp)

print(wfreq.most_common(20))
```

**chunksize** optional parameter in pandas' **read\_json** method reads in 10,000 lines at a time...

then, iterate through each small df.

Memory-efficient!  
This code uses only 290MB of memory!!

# Pandas vs. large data: tips

---

- ▶ "Why and How to Use Pandas with Large (but not big) Data"
  - ◆ <https://towardsdatascience.com/why-and-how-to-use-pandas-with-large-data-9594dda2ea4c>
- 1. Read CSV file data in chunk size
- 2. Filter out unimportant columns in **df** to save memory
- 3. Change **dtypes** for columns
  - ◆ float64 takes up more space than float32.

# Vectorizing and training in chunks

```
from sklearn.naive_bayes import MultinomialNB
from sklearn.feature_extraction.text import HashingVectorizer
import warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)
```

```
filename = 'review_10k.json'
length = 10000
chunk_size = 1000
chunks = length/chunk_size
```

```
df_chunks = pd.read_json(filename, lines=True, chunksize=chunk_size, encoding="utf-8")
```

```
clf = MultinomialNB()
vectorizer = HashingVectorizer(non_negative=True)
```

```
for i, df in enumerate(df_chunks):
    if i < 0.8 * chunks:
        clf.partial_fit(vectorizer.transform(df['text']), df['stars'], classes=[1,2,3,4,5])
    else:
        pred = clf.predict(vectorizer.transform(df['text']))
        print('batch {}, {} accuracy'.format(i, np.mean(pred == df['stars'])))
```

```
batch 8, 0.444 accuracy
batch 9, 0.439 accuracy
```

If vectorizer/ML model depends only on individual row of data, it can be implemented in chunks.  
(Caveat: TF-IDF vectorizer and most ML models can't.)

Hashing vectorizer  
skips the IDF part of TF/IDF,  
can be implemented in chunks!

NB classifier can be  
trained in partial bits!

# Computational efficiency: space vs. time

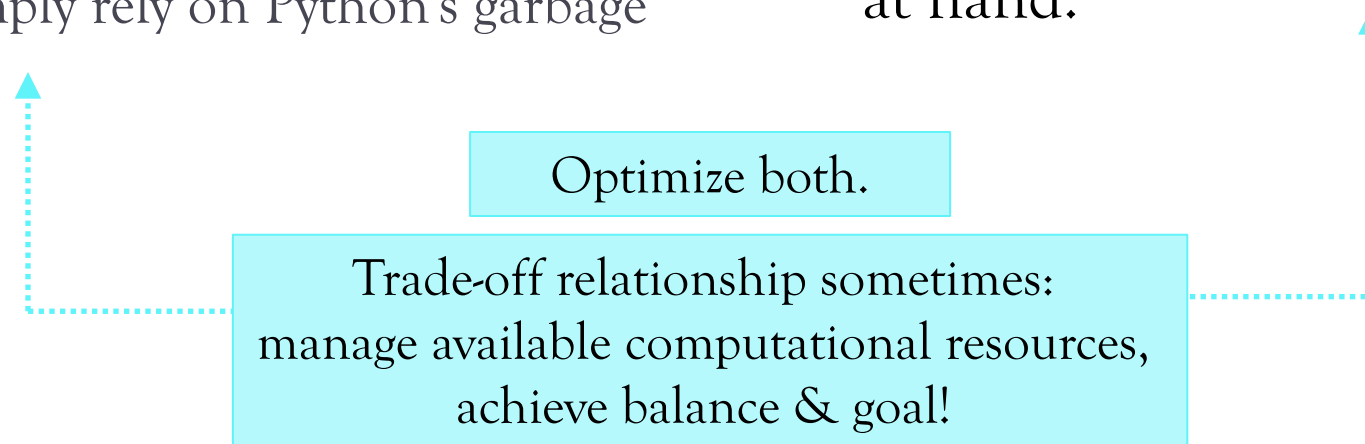
---

## SPACE: memory footprint

- ▶ Do not create duplicate data objects.
- ▶ Avoid creating a data object that does not need to be stored in its entirety.
- ▶ In code, delete large data objects that will no longer be used
  - Do not simply rely on Python's garbage collection

## TIME: processor runtime

- ▶ Avoid duplicating an expensive processing step: process once, store result as an object, then reuse.
- ▶ Use an efficient algorithm.
- ▶ Use the data type optimal for the task at hand.



# Data types and optimization

```
[1]: import nltk
nltk.data.path.append('/zfs2/ling1340-2019s/shared_data/nltk_data')
```

```
[2]: from nltk.corpus import gutenberg
%pprint
gutenberg.fileids()
```

Pretty printing has been turned OFF

```
[2]: ['austen-emma.txt', 'austen-persuasion.txt', 'austen-sense.txt', 'bible-kjv.txt', 'blake-poems.txt', 'bryan
t-stories.txt', 'burgess-busterbrown.txt', 'carroll-alice.txt', 'chesterton-ball.txt', 'chesterton-brown.tx
t', 'chesterton-thursday.txt', 'edgeworth-parents.txt', 'melville-moby_dick.txt', 'milton-paradise.txt', 's
hakespeare-caesar.txt', 'shakespeare-hamlet.txt', 'shakespeare-macbeth.txt', 'whitman-leaves.txt']
```

```
[3]: awords = gutenberg.words('carroll-alice.txt')
print(awords[:100])
print(len(awords))
```

```
['[', 'Alice', '"', 's', 'Adventures', 'in', 'Wonderland', 'by', 'Lewis', 'Carroll', '1865', ']', 'CHAPTE
R', 'I', '.', 'Down', 'the', 'Rabbit', '-', 'Hole', 'Alice', 'was', 'beginning', 'to', 'get', 'very', 'tire
d', 'of', 'sitting', 'by', 'her', 'sister', 'on', 'the', 'bank', ',', 'and', 'of', 'having', 'nothing', 't
o', 'do', ':', 'once', 'or', 'twice', 'she', 'had', 'peeped', 'into', 'the', 'book', 'her', 'sister', 'wa
s', 'reading', ',', 'but', 'it', 'had', 'no', 'pictures', 'or', 'conversations', 'in', 'it', ',', '"', 'an
d', 'what', 'is', 'the', 'use', 'of', 'a', 'book', '"', 'thought', 'Alice', '"', 'without', 'pictures', 'o
r', 'conversation', '"?', 'So', 'she', 'was', 'considering', 'in', 'her', 'own', 'mind', '(', 'as', 'well',
'as', 'she', 'could', ',', '']
34110
```

"Alice in Wonderland",  
34K tokens

```
[4]: enable = open('/zfs2/ling1340-2019s/shared_data/enable1.txt').read().split()
print(enable[:30])
print(len(enable))
```

```
['aa', 'aah', 'aahed', 'aahing', 'aahs', 'aal', 'aalii', 'aaliis', 'aals', 'aardvark', 'aardvarks', 'aardwo
lf', 'aardwolves', 'aargh', 'aarrgh', 'aarrghh', 'aas', 'aasvogel', 'aasvogels', 'ab', 'aba', 'abaca', 'aba
cas', 'abaci', 'aback', 'abacterial', 'abacus', 'abacuses', 'abaft', 'abaka']
172820
```

"enable" word list,  
173K total words

Task: find Alice  
words that are not  
found in enable list

- Try 1: list-comprehend through awords (list), filter against enable (list)

```
[5]: %time notfound = [w for w in awords if w not in enable]
```

- Try 2: same, but filter against enable list as a SET

```
[6]: enable_set = set(enable) # this one is a set data type  
print(len(enable_set))      # same size
```

172820

```
[7]: %time notfound = [w for w in awords if w not in enable_set]
```

- Try 3: compute the SET difference

```
[8]: awords_set = set(awords)  
print(len(awords_set))    # now a set, smaller size
```

3016

```
[9]: %time notfound = awords_set.difference(enable_set)
```

- Try 1: list-comprehend through awords (list), filter against enable (list)

```
[5]: %time notfound = [w for w in awords if w not in enable]
```

```
CPU times: user 39.7 s, sys: 16.2 ms, total: 39.7 s  
Wall time: 39.8 s
```

- Try 2: same, but filter against enable list as a SET

```
[6]: enable_set = set(enable) # this one is a set data type  
print(len(enable_set)) # same size
```

```
172820
```

```
[7]: %time notfound = [w for w in awords if w not in enable_set]
```

```
CPU times: user 20.3 ms, sys: 4 ms, total: 24.3 ms  
Wall time: 24.8 ms
```

much  
faster

Lists as a data type  
are NOT optimized  
for membership  
operations...

but sets are!

- Try 3: compute the SET difference

```
[8]: awords_set = set(awords)  
print(len(awords_set)) # now a set, smaller size
```

```
3016
```

```
[9]: %time notfound = awords_set.difference(enable_set)
```

```
CPU times: user 332 µs, sys: 2 µs, total: 334 µs  
Wall time: 339 µs
```

blazing  
fast

Keep efficiency in  
mind: pick the  
right combination  
of data structure  
and operation



# Algorithmic complexity and the Big O

---

▶ <https://rob-bell.net/2009/06/a-beginners-guide-to-big-o-notation/>

▶ We have a **list of  $n$  items**. Imagine  $n$  is 100, 1000 or even 1 million.

1. Is the first element an even number?

- ◆ Can be implemented in  $O(1)$ : an algorithm that executes in a **constant** time regardless of the size of the input dataset.

2. Does the list contain value 42?

- ◆ Can be implemented in  $O(n)$ : an algorithm whose performance will grow **linearly** in proportion to the size of the input data.

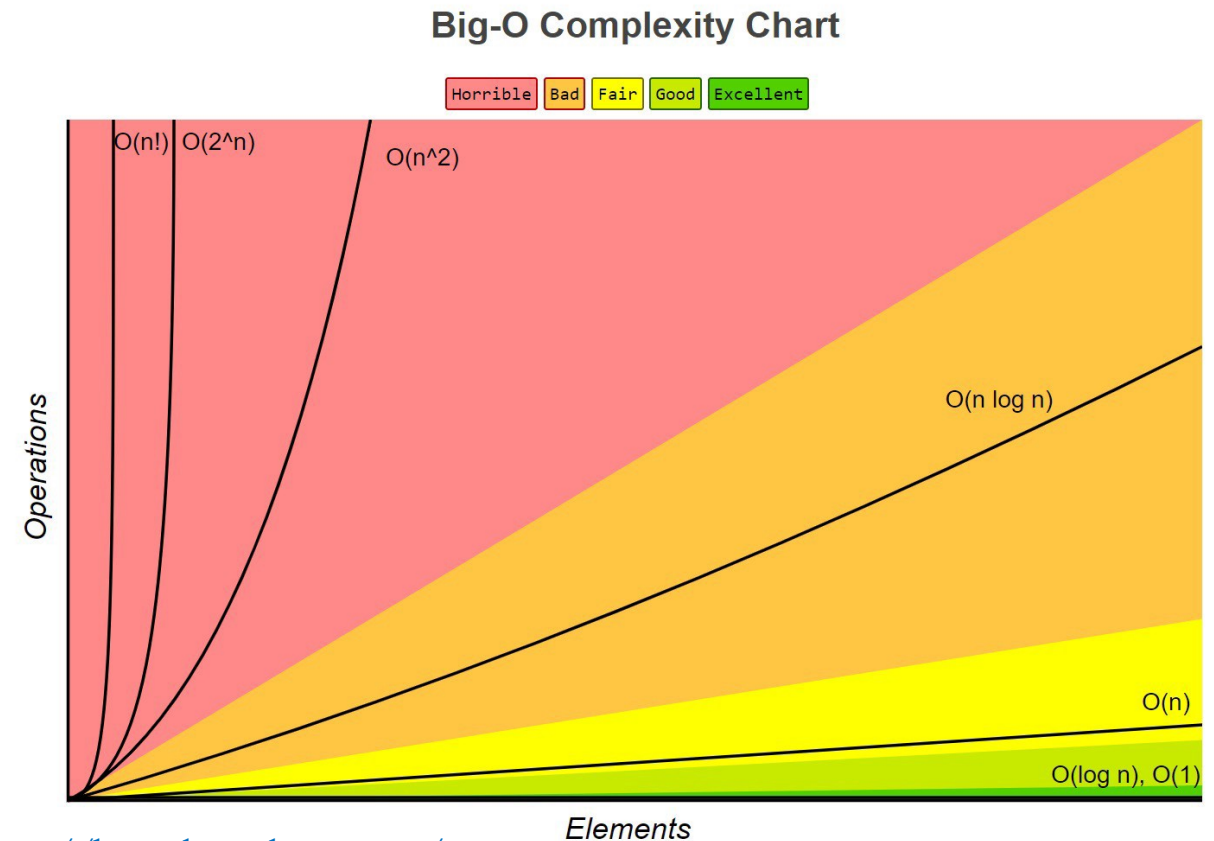
3. Does the list contain duplicate values?

- ◆ Can be implemented in  $O(n^2)$ : an algorithm whose performance is directly proportional to the square of the size of the input data set (**quadratic**).

# Algorithmic complexity and the Big O

- ▶ <https://rob-bell.net/2009/06/a-beginners-guide-to-big-o-notation/>
- ▶ We have a **list of  $n$  items**. Imagine  $n$  is 100, 1000 or even 1 million.

4. Sort the list (ascending or descending)
  - ♦ Can be implemented in  $O(n \log n)$ : an algorithm that executes in loglinear time.
  - ♦ See: <https://brilliant.org/wiki/sorting-algorithms/>



# Algorithmic efficiency: summary

---

- ▶ A problem can be implemented with varying degrees of algorithmic efficiency.
- ▶ A problem comes with its own inherent algorithmic complexity limit.
  - ◆ **Big O notation** is a mathematical notation that encapsulates the relationship between the processing time and the input data size.
  - ◆ Example: the most efficient known sorting algorithm bottoms out at  $O(n \log n)$ .
- ▶ In a nutshell...
  - ◆ Compose the most efficient algorithm that you can.
  - ◆ Understand the relationship between the data size growth and the processing time growth.  $O(n)$  has fair scalability,  $O(n^2)$  becomes intractable.
  - ◆ Efficiency of an algorithm can lead to dramatic runtime difference when dealing with big data.

# Word representations

---

- ▶ We've seen Count Vectorizer and TF-IDF
- ▶ These don't account for:
  - ◆ Word meaning
  - ◆ Word context
- ▶ How can we capture things like:
  - ◆ “apple” and “orange” are fruits
  - ◆  $king - man + woman = queen$

# Word embeddings (simplified)

---

- ▶ An embedding of a word is a vector representing that word
  - ◆ Example 4-dimensional embedding for *dog*:  $[-0.4, 0.37, 0.02, -0.34]$
- ▶ Each dimension of the vector represents some meaning or concept

		Dimensions					
Word vectors	dog	-0.4	0.37	0.02	-0.34	animal	
	cat	-0.15	-0.02	-0.23	-0.23	domesticated	
	lion	0.19	-0.4	0.35	-0.48	pet	
	tiger	-0.08	0.31	0.56	0.07	fluffy	
	elephant	-0.04	-0.09	0.11	-0.06		
	cheetah	0.27	-0.28	-0.2	-0.43		
	monkey	-0.02	-0.67	-0.21	-0.48		
	rabbit	-0.04	-0.3	-0.18	-0.47		
	mouse	0.09	-0.46	-0.35	-0.24		
	rat	0.21	-0.48	-0.56	-0.37		

# Word embeddings (less simple)

---

- ▶ In reality, what each dimension represents is much harder to define
- ▶ Words with similar meanings have similar representations
- ▶ While TF-IDF focuses on word relevance, word embeddings focus on word meaning
- ▶ Many different ways to produce word embeddings...

# Word2Vec

---

- ▶ First major model of word embeddings
  - ◆ Introduced by Google Research in 2013
- ▶ Given some text, for each word, take the word and some context
  - ◆ One way to do this is a *skip-gram model*: Given a “center” word, use a neural network to predict the probabilities that it will be surrounded by some context words
  - ◆ Another way is to train a model to “fill in” the “center” word given some context
    - ◆ “the cat \_ on the mat” > high probability for “sat”
  - ◆ Embeddings are produced by an intermediate “hidden” layer in the network

# Why do this?

- ▶ Vector arithmetic with words!  
(*king* - *man* + *woman* = *queen*)
- ▶ Show relationships between concepts in 2D

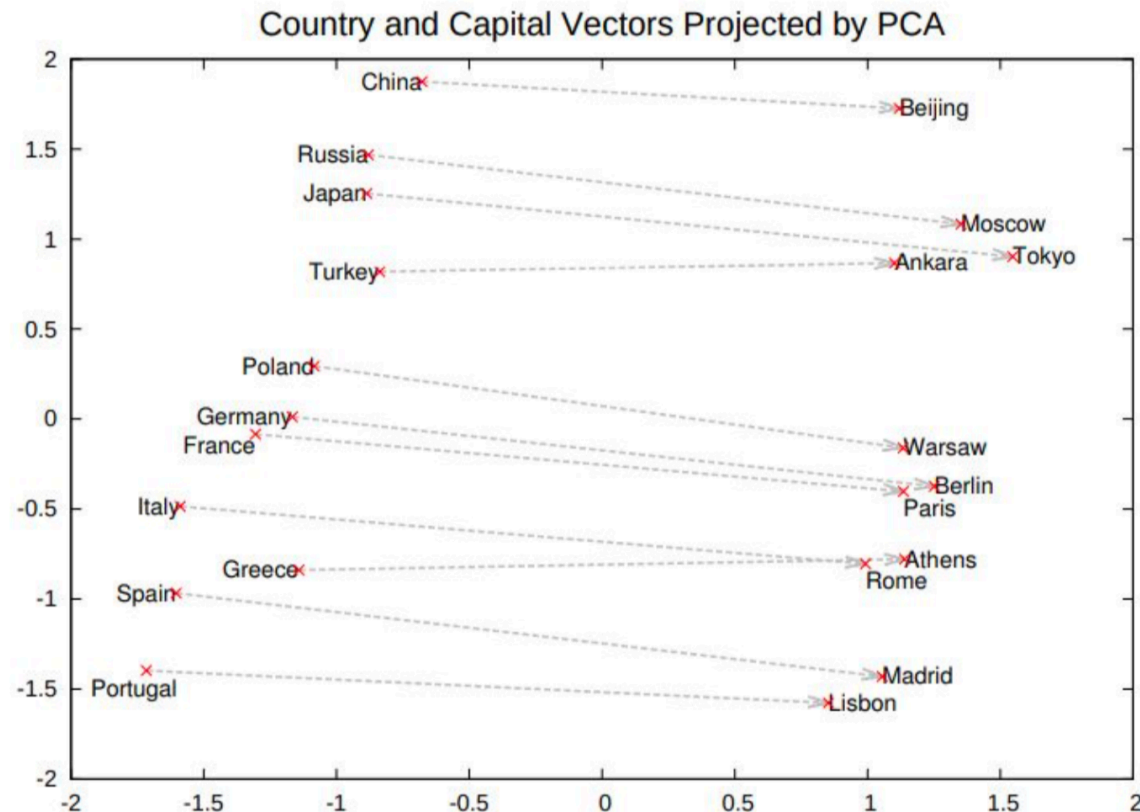


Figure 2: Two-dimensional PCA projection of the 1000-dimensional Skip-gram vectors of countries and their capital cities. The figure illustrates ability of the model to automatically organize concepts and learn implicitly the relationships between them, as during the training we did not provide any supervised information about what a capital city means.



# GloVe: Global Vectors for Word Representation

- ▶ Not training a model to predict stuff; rather, using co-occurrence ratios
  - ◆  $P(k|w)$  = probability that  $k$  appears in the context of  $w$ .
  - ◆ *solid* is related to *ice* but not to *steam*, so  $P(\text{solid}|\text{ice})$  should be high and  $P(\text{solid}|\text{steam})$  low.
  - ◆ So the ratio of  $P(\text{solid}|\text{ice}) / P(\text{solid}|\text{steam})$  is large.
  - ◆ *gas* is related to *steam* but not *ice*, so the ratio of  $P(\text{gas}|\text{ice}) / P(\text{gas}|\text{steam})$  is small.
  - ◆ Words related to both (e.g. *water*) or neither (e.g. *fashion*) should have a ratio close to 1.
  - ◆ The dot product of two vectors equals the log of the words' probability of co-occurrence.

Probability and Ratio	$k = \text{solid}$	$k = \text{gas}$	$k = \text{water}$	$k = \text{fashion}$
$P(k \text{ice})$	$1.9 \times 10^{-4}$	$6.6 \times 10^{-5}$	$3.0 \times 10^{-3}$	$1.7 \times 10^{-5}$
$P(k \text{steam})$	$2.2 \times 10^{-5}$	$7.8 \times 10^{-4}$	$2.2 \times 10^{-3}$	$1.8 \times 10^{-5}$
$P(k \text{ice})/P(k \text{steam})$	8.9	$8.5 \times 10^{-2}$	1.36	0.96

# FastText

---

- ▶ Basically an extension of Word2Vec
  - ◆ Facebook's AI Research lab, 2015
- ▶ Uses character ngrams, e.g. *apple* > *ap*, *app*, *appl*, *apple*, *pp*, *ppl*, *pple*, etc.
- ▶ Benefits:
  - ◆ Can generate word embeddings for new words, as long as it has the right character ngrams
  - ◆ Supposedly better performance than other models for certain downstream tasks
- ▶ Drawbacks:
  - ◆ More hyperparameters to tweak (including ngram sizes)
  - ◆ Requires more resources in both time and space

# More on word embeddings

---

- ▶ Word embeddings (usually pretrained) are the standard features to start with for basically any NLP research task
  - ◆ image captioning, question answering, machine translation, what-have-you
- ▶ Steps for text classification using word embeddings:
  - ◆ Tokenize text
  - ◆ Map each token to its word embedding to create 2D matrix of features
  - ◆ Add more features if desired
  - ◆ Train classifier
- ▶ Fixed-length input models can average our word embeddings to create a single feature vector
- ▶ Variable-length input models (e.g., recurrent neural networks) lose less information

# Wrapping up

---

## ► Homework 4

- ◆ Fun with CRC!
- ◆ Make sure to check out the list of tips at the bottom of the page...
- ◆ As a class we have 10000 Service Units, which is 10k hours of computing time (not that much).
  - ◆ Don't let your jobs run infinitely! Test on small samples of data and make sure they work before you run large jobs.

## ► Next class

- ◆ Clustering, HW review